

Final 37810

Shiting Zhu, Lijing Wang, Jingting Li

November 1, 2015

Question 1

Metropolis Hastings

1. The algorithm is Metropolis-Hastings. We first choose the new phi based on the proposal function.

- Step 0: Pick a start point ϕ_0 , which is a random number from uniform distribution (0,1)
- Step 1: Pick a new candidate(ϕ_{new}) from the jumping distribution, which is $\text{Beta}(c\phi_{old}, c(1-\phi_{old}))$. Notice the ϕ_{new} cannot be either 0 or 1, otherwise it will loop in 0 or 1 forever. For this step we used a function called phi in the r code.
- Step 2: Calculate the acceptance probability. Since the jumping distribution is not symmetric in this question, so

$$\frac{P(\phi_{new})}{P(\phi_n)} = \frac{dbeta(\phi_{new}, 6, 4)}{dbeta(\phi_n, 6, 4)}$$

$$\frac{Q(\phi_{new} * |\phi_n)}{Q(\phi_n|\phi_{new})} = \frac{dbeta(\phi_{new}, c\phi_n, c(1-\phi_n))}{dbeta(\phi_n, c\phi_{new}, c\phi_{new})}$$

$$A(\phi_n - > \phi_{new}) = \min(1, \frac{\frac{P(\phi_{new})}{P(\phi_n)}}{\frac{Q(\phi_{new} * |\phi_n)}{Q(\phi_n|\phi_{new})}})$$

- Step 3: Sample a number from uniform(0,1). If it is more than $A(\phi_n - > \phi_{new})$, we accept ϕ_{new} and set $\phi_{n+1} = \phi_{new}$. Otherwise, $\phi_{n+1} = \phi_n$.
- Step 4 : Repeat Step 1 to Step 3 until we get length n chain.

```
library(coda)
set.seed(123)

Alpha <- 6
Beta <- 4

# Use the proposal function to get new phi, noticeing phi cannot be 0 or 1
phi <- function(c,oldphi){
  newphi = 0
  # Making sure new phi is neither 0 nor 1
  while (newphi == 0 || newphi == 1){
    newphi = rbeta(1,c*oldphi,c*(1-oldphi))
  }
  return(newphi)
}
```

```

# Using similar strcuture of assignment 3 to build MCMC chain
run_metropolis_MCMC <- function(startvalue,c,iterations){
  # set chain
  chain = rep(0, iterations+1)
  chain[1] = startvalue
  for (i in 1:iterations){
    phi = phi(c,chain[i])
    posterior = dbeta(phi,Alpha,Beta)/dbeta(chain[i],Alpha,Beta)
    proposal = dbeta(phi,c*chain[i],c*(1-chain[i]))/dbeta(chain[i],c*phi,c*(1-phi))
    probab = min(1,posterior/proposal)
    # accept new value if random number uniform (0,1) is less than
    # acceptance probability
    if (runif(1) < probab){
      chain[i+1] = phi
      # reject new value if random number uniform (0,1) is greater
      # or equal than acceptance probability
    }else{
      chain[i+1] = chain[i]
    }
  }
  return(chain)
}

startvalue = runif(1)
chain = run_metropolis_MCMC(startvalue,1, 10000)
acceptance = 1-mean(duplicated(chain))

```

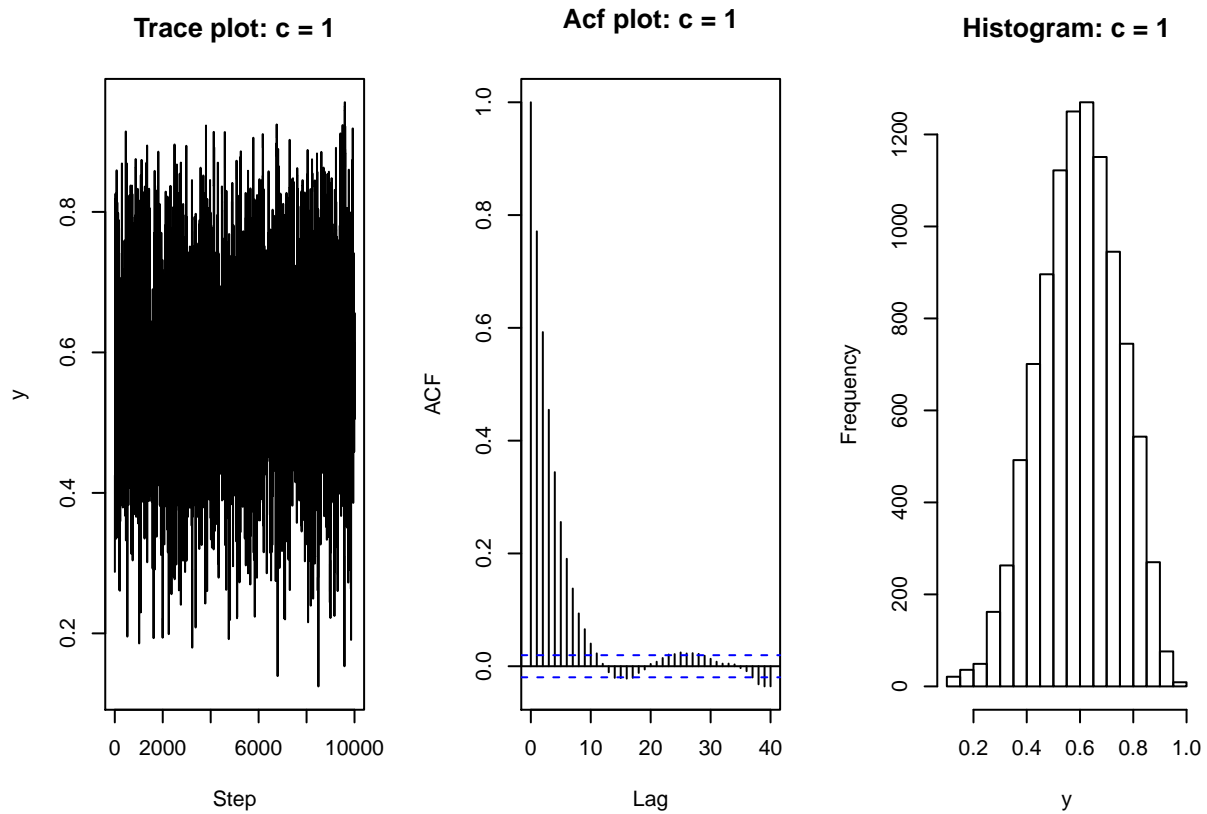
2. Based on the plots, we think the performance of the sampler when $c = 1$ is ok. The traceplot shows the movement of the MCMC chain values. The autocorrection plot is good and not much thinning is needed. The histograms of the MCMC chain and target distribution $\text{Beta}(6,4)$ are pretty close. The Kolmogorov-Smirnov Statistic is showing good result and supporting our conclusions.

```

par(mfrow=c(1,3)) #1 row, 3 columns

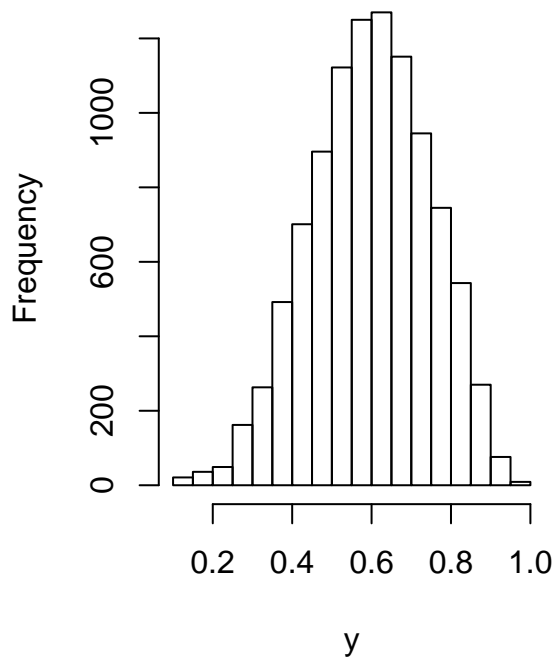
traceplot(as.mcmc(chain), type="l", main = "Trace plot: c = 1", xlab="Step", ylab="y")
acf(chain, main = "Acf plot: c = 1")
hist(chain, main = "Histogram: c = 1", xlab="y")

```




```
target = rbeta(10001,6,4)
par(mfrow=c(1,2))
hist(chain, main = "Histogram: c = 1", xlab="y")
hist(target, main = "Histogram: Beta(6,4)", xlab="y")
```

Histogram: $c = 1$



Histogram: Beta(6,4)



```
ks.test(chain,"pbeta",6,4)
```

```
## Warning in ks.test(chain, "pbeta", 6, 4): ties should not be present for
## the Kolmogorov-Smirnov test
```

```
##
## One-sample Kolmogorov-Smirnov test
##
## data: chain
## D = 0.022219, p-value = 0.0001029
## alternative hypothesis: two-sided
```

3. From the result we can see when $c = 2.5$ gives the best result. Based on the autocorrelation plots, $c = 0.1$ is pretty bad since it requires a lot of burn-in. $c = 10$ and $c = 2.5$ have very close autocorrelation plots but when $c = 2.5$ it is a little bit better. $c = 0.1$ histogram is very different than our target distribution while $c = 2.5$ and $c = 10$ histograms are very close to our target distribution histogram. The Kolmogorov-Smirnov Statistic result is consistent with our conclusion.

```
Cs <- c(0.1, 2.5, 10)
# calculate MCMC chain for all the c
result <- sapply(Cs, run_metropolis_MCMC, startvalue = startvalue, iterations = 10000)

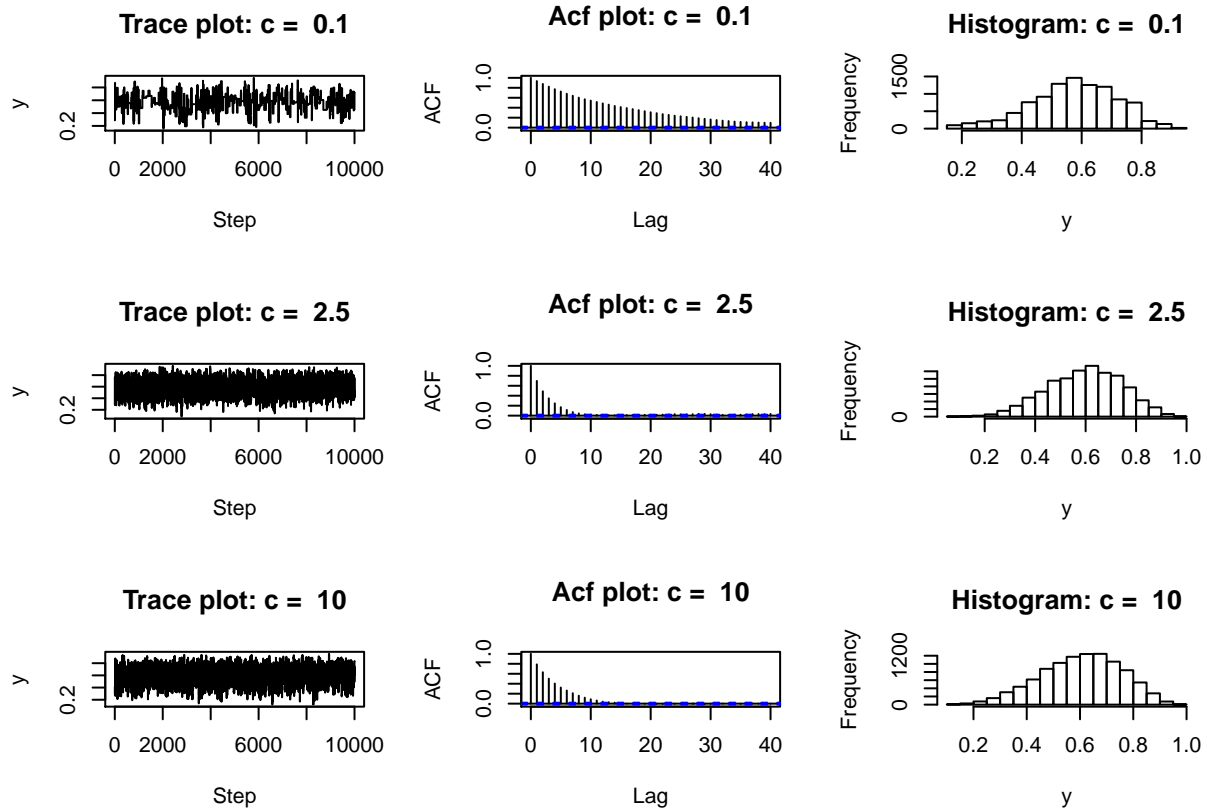
# set a graph function which will automatically have the trace plot,
# cf plot and histograms of all the c in vector Cs
graph <- function(Cs){
  n = length(Cs)
  par(mfrow = c(n, 3))
  for( i in 1:n ) {
```

```

traceplot(as.mcmc(result[,i]), type="l", main = paste("Trace plot: c = ", Cs[i]), xlab="Step", ylab="y")
acf(result[,i], main = paste("Acf plot: c = ", Cs[i]))
hist(result[,i], main = paste("Histogram: c = ", Cs[i]), xlab="y")
}
}

```

```
graph(Cs)
```

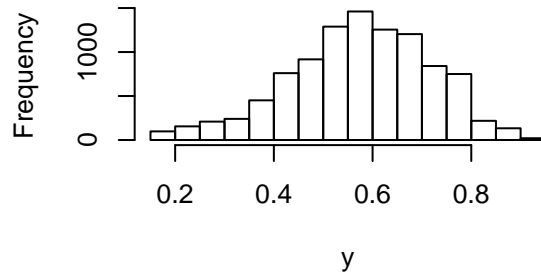


```

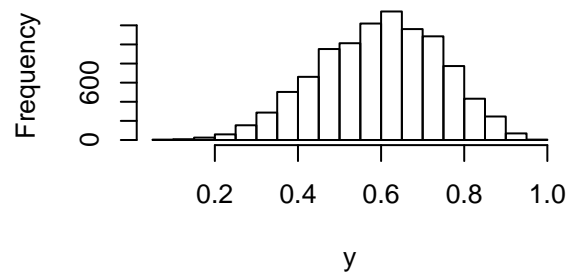
par(mfrow=c(2,2))
for( i in 1:length(Cs) ) {
hist(result[,i], main = paste("Histogram: c = ", Cs[i]), xlab="y")
}
hist(target, main = "Histogram: Beta(6,4)", xlab="y")

```

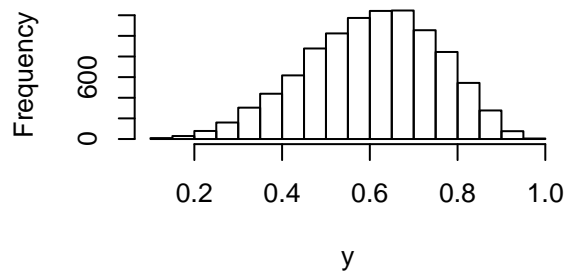
Histogram: c = 0.1



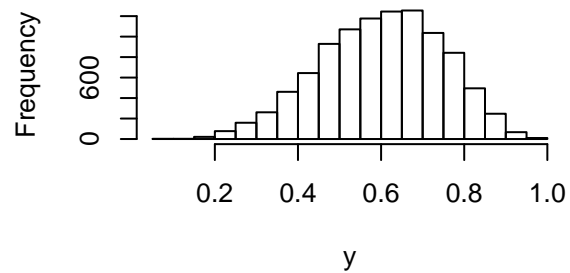
Histogram: c = 2.5



Histogram: c = 10



Histogram: Beta(6,4)



```
ks.test(result[,1],"pbeta",6,4)
```

```
##
## One-sample Kolmogorov-Smirnov test
##
## data: result[, 1]
## D = 0.092739, p-value < 2.2e-16
## alternative hypothesis: two-sided
```

```
ks.test(result[,2],"pbeta",6,4)
```

```
##
## One-sample Kolmogorov-Smirnov test
##
## data: result[, 2]
## D = 0.017795, p-value = 0.003549
## alternative hypothesis: two-sided
```

```
ks.test(result[,3],"pbeta",6,4)
```

```
##
## One-sample Kolmogorov-Smirnov test
##
## data: result[, 3]
## D = 0.015745, p-value = 0.01405
## alternative hypothesis: two-sided
```

Question 2

Gibbs Sampling

```
Gibbs_<-function(x0,y0,iterations,B=5,burnIn=0)
## 5 parameters are needed in this function, with B=5 and burnIn=0 by default
{
  if(x0>0&& x0<B&&y0>0&&y0<B)
## to check whether the starting values are in the domain.
  {
    x<-c(x0,rep(NA,iterations-1))
## Initialize the Markov chain
    y<-c(y0,rep(NA,iterations-1))
## Initialize the Markov chain
    for(i in 1:(iterations-1))
    {
      x[i+1]<-(-log(1-runif(1)*(1-exp(-y[i]*B)))/y[i])
## use inverse transform sampling to draw sample from conditional distribution
## p(x^{i+1}|y^i)
      y[i+1]<-(-log(1-runif(1)*(1-exp(-x[i+1]*B)))/x[i+1])
## use inverse transform sampling to draw sample from conditional distribution
## p(y^{i+1}|x^{i+1})
    }
    if(burnIn>0)
    {
      x<-x[-(1:burnIn)]
      y<-y[-(1:burnIn)]
      print(length(x))
## discard the first bunch of draws for the burn-in process
    }
    return(data.frame(x,y))
  }
  else
    stop("Initial values incorrect")
## print the information for incorrect starting values
}
```

To estimate the marginal distribution generated from the conditional distributions using Gibbs sampling, we first pick up the starting values for X and Y . Denote them as x_0 and y_0 . `iterations` stands for the number of draws we want to get from Gibbs sampling. B is the number given in the exercise, which is 5 here. `burnIn` is the number of draws we want to discard for the burn-in process, the default for `burnIn` is 0.

In this case, both the conditional distribution of $X|Y$ and $Y|X$ are truncated exponential distribution. The domains of both conditional pdf's are $[0, B]$. So the starting values should satisfy $0 < x_0 < B$ and $0 < y_0 < B$. Therefore I put a restriction `if(x0>0&&x0<B&&y0>0&&y0<B)` on the input of starting values to check if they satisfy the requirement.

Since

$$p(x|y) \propto ye^{-yx}, 0 < x < B$$

$$p(y|x) \propto xe^{-yx}, 0 < y < B$$

Consider $p(x|y)$ only. To get samples using inverse transform sampling, first we need to figure out the normalizing constant for the conditional pdf and then the inverse function of cdf.

The normalizing constant can be calculated using the formula $c = \frac{1}{\int f(x|y) dx}$, if $p(x|y) \propto f(x|y)$. In this case, $f(x|y) = ye^{-yx}$, $0 < x < B$.
Therefore

$$c = \frac{1}{\int_0^B ye^{-yx} dx} = \frac{1}{1 - e^{-By}}$$

Then we start to calculate the cdf $H(x|y)$ for $X|Y$.

We have

$$p(x|y) = \frac{ye^{-yx}}{1 - e^{-By}}, 0 < x < B$$

Using the formula $F(x) = \int_{-\infty}^x p(z) dz$, we get the cdf of $X|Y$:

$$F(x|y) = \frac{1 - e^{-xy}}{1 - e^{-By}}, 0 < x < B$$

Then we can write down the inverse function of $F(x|y)$, Denote as $F^{-1}(u|y)$, where $0 \leq u \leq 1$.

We have:

$$F^{-1}(u|y) = -\frac{\log(1 - u(1 - e^{-By}))}{y}, 0 \leq u \leq 1$$

According to inverse transform sampling, to draw a sample from $p(x|y)$, we first generate u from $Unif[0, 1]$, then $x = F^{-1}(u|y)$ is from the conditional distribution of $X|Y$.

We draw samples from $p(y|x)$ using the same method.

Now we start generating samples from $p(x, y)$.

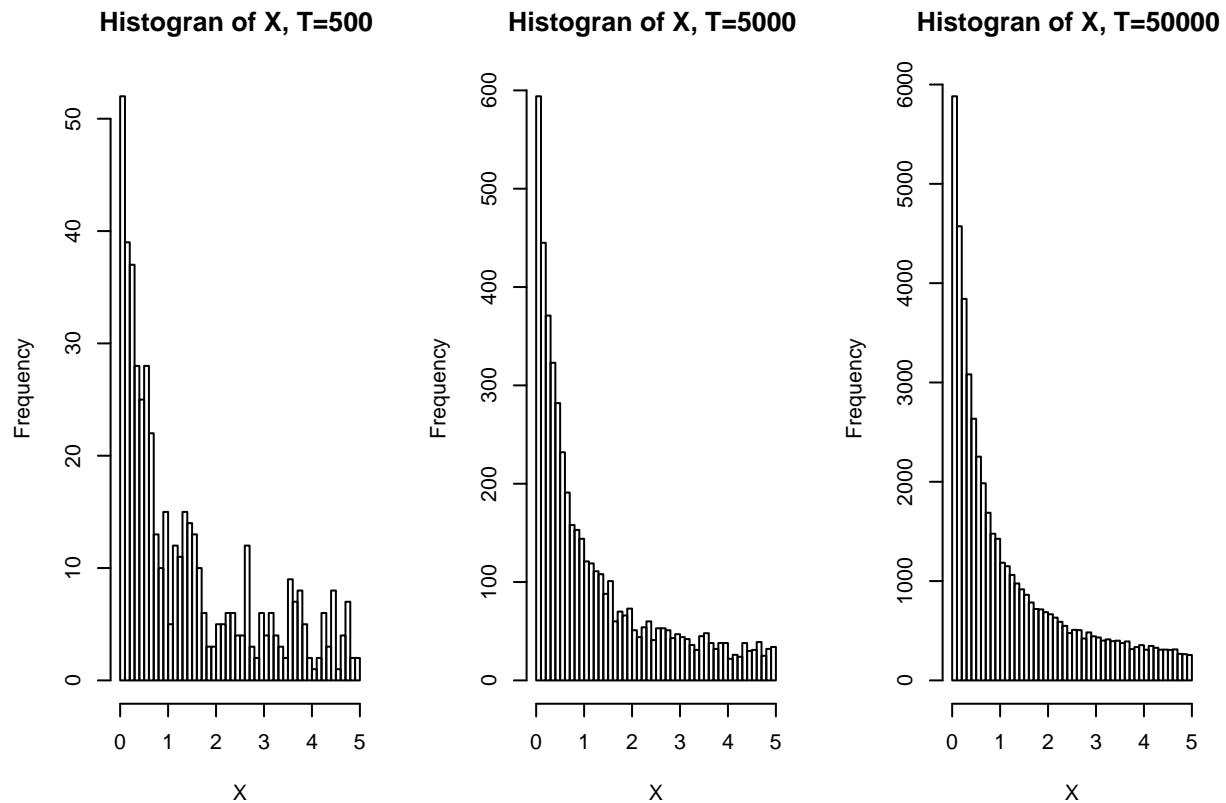
Here I start Gibbs sampling with x_0 and y_0 , draw a value $x^{(1)}$ from the full conditional $p(x|y_0)$ using inverse transform sampling. Then use the updated $x^{(1)}$ to draw a sample from $p(y|x^{(1)})$.

To get more samples using Gibbs sampling, continually use the most updated values of x and y when generating samples from conditional distribution. To be more specific, generate $x^{(i+1)}$ from $p(x|y^{(i)})$ and $y^{(i+1)}$ from $p(y|x^{(i+1)})$. Repeat it for $n = \text{iterations}$ times.

For the burn-in process, discard the first $m = \text{burnIn}$ samples from the Markov chain to reduce the influence of starting values on the Markov chain.

For $B = 5$ and sample size $T = 500, 5000, 50000$, to see the histogram of values for x generated from Gibbs sampling, We have:

```
par(mfrow=c(1,3))
x<-Gibbs_(1,2,50000,B=5,burnIn=0)
hist(x[1:500,1],breaks=60, main="Histogram of X, T=500 ",xlab="X")
hist(x[1:5000,1],breaks=60, main="Histogram of X, T=5000 ",xlab="X")
hist(x[1:50000,1],breaks=60, main="Histogram of X, T=50000 ",xlab="X")
```

Here we picked up the starting values arbitrarily.

The plots above shows that as T increases, the histogram tend to be more smooth.

Use the mean of the generated samples to estimate the expectation of X .

```
set.seed(123)
mean_500<-c()
mean_5000<-c()
mean_50000<-c()
for(i in 1:10)
{
  x<-Gibbs_(1,2,50000,B=5,burnIn=0)
  mean_500<-c(mean_500,mean(x[1:500,1]))
  mean_5000<-c(mean_5000,mean(x[1:5000,1]))
  mean_50000<-c(mean_50000,mean(x[1:50000,1]))
}
mean_500

## [1] 1.323771 1.392163 1.425209 1.300402 1.319226 1.200004 1.389191
## [8] 1.305194 1.360014 1.357299

mean_5000

## [1] 1.266212 1.263367 1.264459 1.280466 1.269729 1.280995 1.252972
## [8] 1.279271 1.254663 1.296133
```

```
mean_50000
```

```
## [1] 1.259416 1.272679 1.266894 1.275371 1.269562 1.270537 1.267437  
## [8] 1.267778 1.261333 1.262746
```

```
c(var(mean_500),var(mean_5000),var(mean_50000))
```

```
## [1] 0.0040109521 0.0001785322 0.0000253901
```

As we can see, the mean is around 1.26, as sample size T increases, the variance of the estimator also decreases.